

Homework 3

Due: 11.59pm on Sunday, February 26

Submission instructions:

- Submit one write-up per group on gradescope.com. Please do not bring printouts of your solutions to the classroom.

Question 1

In this homework, you will be analyzing the data coming from KDD Cup 2009: Customer relationship prediction. The KDD Cup is a annual data mining competition and the data were given by French Telecom company Orange. There are 3 (binary) response variables to explore:

- churn: predicting the propensity of customers to switch provider
- appetency: predicting the propensity of customers to buy new products or services
- up-selling: predicting the propensity of customers to buy upgrades or add-ons proposed to them to make the sale more profitable

We are using the small dataset consists of 50,000 observations and 230 explanatory variables (the large one comes with 15,000 explanatory variables). The first 190 variables are numerical and the last 40 are categorical. Due to the nature of competition and privacy concerns, none of the explanatory variables come with meaningful names, are the categorical variables are encoded anonymously as well.

1. Getting the Data

Data are located in the repository here:

https://github.com/ChicagoBoothML/MLClassData/tree/master/KDDCup2009_Customer_relationship

The “orange_small_train.data.gz” contains the explanatory variables. We load it in R as X. There are three txt files storing the response variables, we code them as Y_churn, Y_appetency and Y_upselling.

The following code will help you load it into R.

```
if("R.utils" %in% rownames(installed.packages()) == FALSE) {install.packages("R.utils")}
if("data.table" %in% rownames(installed.packages()) == FALSE) {install.packages("data.table")}
library("R.utils")
library("data.table")

gitURL="https://github.com/ChicagoBoothML/MLClassData/raw/master/KDDCup2009_Customer_relationship/";
DownloadFileList=c("orange_small_train.data.gz","orange_small_train_appetency.labels.txt",
                  "orange_small_train_churn.labels.txt","orange_small_train_upselling.labels.txt")
LoadFileList=c("orange_small_train.data","orange_small_train_appetency.labels.txt",
              "orange_small_train_churn.labels.txt","orange_small_train_upselling.labels.txt")

for (i in 1:length(LoadFileList)){
  if (!file.exists(LoadFileList[[i]])){
    if (LoadFileList[[i]]!=DownloadFileList[[i]]) {
      download.file(paste(gitURL,DownloadFileList[[i]],sep=""),
                    destfile=DownloadFileList[[i]])
      gunzip(DownloadFileList[[i]])
    }
  }
}
```

```

    }else{
        download.file(paste(gitURL,DownloadFileList[[i]],sep=""),
                      destfile=DownloadFileList[[i]])}
}

na_strings <- c(' ',
  'na', 'n.a', 'n.a.',
  'nan', 'n.a.n', 'n.a.n.',
  'NA', 'N.A', 'N.A.',
  'NaN', 'N.a.N', 'N.a.N.',
  'NAN', 'N.A.N', 'N.A.N.',
  'nil', 'Nil', 'NIL',
  'null', 'Null', 'NULL')

X=as.data.table(read.table('orange_small_train.data',header=TRUE,
                          sep='\t', stringsAsFactors=TRUE, na.strings=na_strings))
Y_churn      =read.table("orange_small_train_churn.labels.txt", quote="\"")
Y_appetency=read.table("orange_small_train_appetency.labels.txt", quote="\"")
Y_upselling=read.table("orange_small_train_upselling.labels.txt", quote="\"")

```

2. Cleaning the Data

The data we have is challenging for a number of reasons. We list two major ones below:

- Missing values
 - Most columns contain missing values. Some columns contain only missing values, remove such columns. Some columns contain mostly missing values, remove them as well. There might be columns containing only 1 value other than missing, remove them as well. For numeric variables, there are a number of ways to deal with missing values. The simplest one is to simply impute the missing value with the mean of observed values. That is, if a particular value is missing in a column just simply substitute NA with the mean of the column. For categorical variables, the easiest way to deal with missing values is to treat them as if they are one of the levels. For example, if a feature takes on values a, b, and c, then the missing value NA can be treated simply as value d.
- Categorical variables with a large number of levels
 - Some categorical features have a large number of different levels. For example, if you try to fit a decision tree using the features, the algorithms will need to make a large number of splits. One way around this is to combine or aggregate levels for which there are a few observations. One suggestions would be as follows: create a new level “low”, which combines all existing levels for which we have 1~249 observations; create a level “medium” that aggregates existing levels for which there are 250~499 observations; create level “high” that aggregates all existing leveles for which there are 500~999 observations; and keep all other levels as they are. There will be a lot of explanatory variables than a simple linear model can handle, and most of those variables do not have any explanatory power.

Coding Hints:

1. How to write a loop over columns of a data.frame (say, X here)?

```

for (i in names(X)) {
  CurrentColumn = X[[i]]
  CurrentColumnVariableName = i
  # Then you do the computation on CurrentColumn, using
  # function is.na, and save the result
  cat(i, mean(is.na(CurrentColumn)), "\n")
}

```

```
}
```

2. How to drop columns of a data.frame indexed by a list?

```
ExcludeVars = c("Var1", "Var2", "Var100") #for example
idx = !(names(X) %in% ExcludeVars)
XS = X[, !(names(X) %in% ExcludeVars), with = FALSE]
```

3. How to convert missing values into factors?

```
i = "Var208" #for example

CurrentColumn = XS[[i]] #Extraction of column
idx = is.na(CurrentColumn) #Locate the NAs
CurrentColumn = as.character(CurrentColumn) #Convert from factor to characters
CurrentColumn[idx] = paste(i, "_NA", sep = "") #Add the new NA level strings
CurrentColumn = as.factor(CurrentColumn) #Convert back to factors
XS[[i]] = CurrentColumn #Plug-back to the data.frame
```

4. How to aggregate a number of factors into new factors?

```
Thres_Low = 249
Thres_Medium = 499
Thres_High = 999
i = "Var220" #for example, this one has 4291 levels
CurrentColumn = XS[[i]] #Extraction of column

CurrentColumn_Table = table(CurrentColumn) #Tabulate the frequency
levels(CurrentColumn)[CurrentColumn_Table <= Thres_Low] = paste(i,
  "_Low", sep = "")
CurrentColumn_Table = table(CurrentColumn)
levels(CurrentColumn)[CurrentColumn_Table > Thres_Low & CurrentColumn_Table <=
  Thres_Medium] = paste(i, "_Medium", sep = "")
CurrentColumn_Table = table(CurrentColumn)
levels(CurrentColumn)[CurrentColumn_Table > Thres_Medium & CurrentColumn_Table <=
  Thres_High] = paste(i, "_High", sep = "")

XS[[i]] = CurrentColumn #Plug-back to the data.frame
# We got 9 levels after cleaning
```

Follow the procedures/suggestions to clean the data, and provide summary graphs/tables if necessary. You could use your own procedures as well, just state the rationals. Report the columns you removed in a nice way (a figure would be perfect). After you have done cleaning the data, please split them into two parts: train + validation set (~80% of total observations) and test set (~20%).

3. Feature/Variable Selection

After cleaning of data, (hopefully) there would still be a bunch of variables remain, presenting the challenge of large p in machine learning. A proportion of these remaining variables are not important at all and you may want to exclude them in the model. Pick one response variable Y (churn,appetency,upselling) here, and stick to that for the rest of the homework. Here are some ideas for you to try, for the Y you picked, using the training + validation set only:

- Fit a classifier for Y using only one explanatory variable, rank the variables by performance and exclude the worst ones.
 - This might seems intuitive and fast, it does ignore the effect of interactions between variables.

- Fit a classifier for Y using pairs of explanatory variables, and exclude variables that are not among the top performing pairs.
 - Caution: Although it takes care of pair-wise interaction, this takes a long time to compute (p^2 versus p), you may want to do the experiment on a smaller set of observations (≤ 1000) first, or do the experiment on a smaller set of variables.
- Fit a random forest model for Y and keep the important variables.
 - Caution: Again, start with a smaller data set when running the experiment.

You only need to do one of them, or come up with your own idea of variable selection. Summarize what variables you keep.

4. Training the Models

Now you have a relatively clean data with a subset of variables, fit a final model using the train + validation set by random forest or boosting. Remember to use cross-validation or out-of-bag error estimate to select tuning parameters. Report the performance of your model in train + validation set.

5. Final Evaluation

Use your final model to make predictions on the test set, report the performance.

(6. OPTIONAL BONUS QUESTION)

In the Feature/Variable Selection step, you pick the important features by running models with only one response variable. If you run the step with each of the response variable, and combine the features from 3 Y 's, would the in-sample performance increase? Does it improve the performance in the test data? Run the experiment and check.

Question 2

Smart health and fitness monitoring devices (for example, created by FitBit, Nike, Adidas, Misfit Shine, ...) are very popular. They are also useful for analyzing person's daily physical activities and recommending health-enhancing exercises. At their core, how do these devices work? Before analyzing human activities, the first thing the device needs to do is to recognize which activities are being performed in the first place.

This experiment¹ monitors people carrying Samsung Galaxy smartphones. Measurements are collected using the phones' accelerometers and gyroscopes, while subjects perform 6 different activities:

- WALKING
- WALKING_UPSTAIRS
- WALKING_DOWNSTAIRS
- SITTING
- STANDING
- LAYING

Your task is to classify person's activity based on the phones recordings.

We have preprocessed data for you. To obtain the data in R, run:

```
download.file(  
  paste("https://raw.githubusercontent.com/ChicagoBoothML/MLClassData/master/",  
        "HumanActivityRecognitionUsingSmartphones/ParseData.R", sep=""),  
  "ParseData.R")  
data <- parse_human_activity_recog_data()
```

The preprocessed data has been cleaned and normalized, so you are not expected to do any further cleaning.

Training data is located in

```
data$X_train  
data$y_train
```

while the test data is located in

```
data$X_test  
data$y_test
```

Students using Python can use the script ParseData.Py located at

<https://raw.githubusercontent.com/ChicagoBoothML/MLClassData/master/HumanActivityRecognitionUsingSmartphones/ParseData.py>

¹<http://archive.ics.uci.edu/ml/datasets/Human+Activity+Recognition+Using+Smartphones>

Your tasks:

1. Build a Neural Network model to classify the 6 activity patterns and report your Accuracy on the Test set
2. Build a tree-based model (or a few) to do the same thing. Compare accuracy on the test set with that of Neural Networks.

Note: We do not expect you to spent too much time on this exercise. Simply try to build a few models using the h2o package to see how the number of hidden units and number of epochs affect the performance. You can leave all the other parameters at their default values. If you have time, you could also try regularization techniques based on dropout or l1/l2 penalization. Same for boosting or random forests.

Data and additional information on the experiment can also be obtained from our git repository ² or from the UCI repository³.

²<https://github.com/ChicagoBoothML/MLClassData/tree/master/HumanActivityRecognitionUsingSmartphones>

³<http://archive.ics.uci.edu/ml/datasets/Human+Activity+Recognition+Using+Smartphones>