

# BUS 41204 Review Session 2

Examples using k-NN and Trees

Siying Cao

`siyingc@uchicago.edu`

01/14/2017

# Plan

- ▶ Example using beer review data to
  - ▶ solve parts of Question 2 on hw1
  - ▶ run regression trees with `rpart` and `tree` package
- ▶ k-NN and regression tree when  $x$  is multivariate
- ▶ Classification using k-NN

## Solve question 2 with beer review data

Load packages

```
library(data.table)
library(rpart)
library(rpart.plot)
library(tree)
library(kknn)
```

Read in data

```
setwd("C:/Users/siying/Box Sync/TA/TA BUS41204/my stuff")
dt <- fread("beer_review.csv", header=TRUE,
            stringsAsFactors=TRUE)
names(dt)[2:dim(dt)[2]]
```

```
## [1] "beer/ABV"          "beer/beerId"      "beer/brewerId"
## [4] "beer/name"         "beer/style"       "review/appearan
## [7] "review/aroma"      "review/overall"   "review/palate"
## [10] "review/taste"      "review/timeUnix"  "user/ageInSecon
## [13] "user/birthdayRaw"  "user/birthdayUnix" "user/gender"
## [16] "user/profileName"
```

# Prepare dataset with intended variables

- ▶  $y = score$
- ▶ Dataset with two input variables: *ABV* and *Age*.

```
dt2 <- dt[,.(`review/overall`, `beer/ABV`, `user/ageInSeconds`)]  
# remove rows with missing value  
dt2 <- dt2[complete.cases(dt2*0),]  
n <- dim(dt2)[1]  
colnames(dt2) <- c("score", "abv", "age")  
setkey(dt2, abv, age)
```

- ▶ Dataset with only one input: *ABV*

```
dt1 <- dt2[,.(score,abv)]
```

# Summarize the data

```
summary(dt2)
```

##	score	abv	age
##	Min. :1.000	Min. : 0.500	Min. :7.034e+08
##	1st Qu.:3.500	1st Qu.: 5.500	1st Qu.:9.783e+08
##	Median :4.000	Median : 7.000	Median :1.100e+09
##	Mean :3.911	Mean : 7.451	Mean :1.175e+09
##	3rd Qu.:4.500	3rd Qu.: 9.400	3rd Qu.:1.276e+09
##	Max. :5.000	Max. :39.440	Max. :3.627e+09

# Split the data

... 75% training and 25% test

```
set.seed(123)


n_train <- 7000 # approximately 75% of data
tr_idx <- sample(1:n, n_train)
train <- dt1[tr_idx,]
test <- dt1[-tr_idx,]
```

# Fit linear regression model

```
linear <- lm(score~abv, data=train)
```

Create scatter plot and best linear fit

```
plot(dt1$abv, dt1$score, cex=.5, pch=16)  
yhat=predict(linear, dt1)  
lines(dt1$abv, yhat, col="red", lwd=3)
```



# k-NN using cross-validation

Download cv utilities

```
download.file("https://raw.githubusercontent.com/ChicagoBoothML/  
source("docv.R") #this has docvknn used below
```

Do 5-fold cv twice, 10-fold cv once

```
set.seed(123)  
k_vec <- 2:100 # vector of k to loop over  
cv1 <- docvknn(matrix(dt1$abv,ncol=1), dt1$score, k_vec,  
                nfold=5)
```

```
## in docv: nset,n,nfold: 99 10479 5  
## on fold: 1 , range: 1 : 2096  
## on fold: 2 , range: 2097 : 4192  
## on fold: 3 , range: 4193 : 6288  
## on fold: 4 , range: 6289 : 8384  
## on fold: 5 , range: 8385 : 10479
```



```
cv2 <- docvknn(matrix(dt1$abv,ncol=1), dt1$score, k_vec,  
                nfold=5)
```

```
## in docv: nset,n,nfold: 99 10479 5  
## on fold: 1 , range: 1 : 2096  
## on fold: 2 , range: 2097 : 4192  
## on fold: 3 , range: 4193 : 6288  
## on fold: 4 , range: 6289 : 8384  
## on fold: 5 , range: 8385 : 10479
```

```
cv3 <- docvknn(matrix(dt1$abv,ncol=1), dt1$score, k_vec,  
                 nfold=10)
```

```
## in docv: nset,n,nfold: 99 10479 10  
## on fold: 1 , range: 1 : 1048  
## on fold: 2 , range: 1049 : 2096  
## on fold: 3 , range: 2097 : 3144  
## on fold: 4 , range: 3145 : 4192  
## on fold: 5 , range: 4193 : 5240  
## on fold: 6 , range: 5241 : 6288  
## on fold: 7 , range: 6289 : 7336  
## on fold: 8 , range: 7337 : 8384  
## on fold: 9 , range: 8385 : 9432  
## on fold: 10 , range: 9433 : 10479
```

## k-NN using cross-validation

Compute MSE for each cv (at the vector of k's)

```
cv1=cv1/n  
cv2=cv2/n  
cv3=cv3/n
```


## (Optional) Run cv multiple times using for loop

```
set.seed(123)
cvmean=rep(0,length(k_vec))
ndocv=50
cvmat=matrix(0, length(k_vec), ndocv) # matrix to store results
for (i in 1:ndocv){
  cvtemp=docvknn(matrix(dt1$abv,ncol=1), dt1$score,
                  k_vec, nfold=10)
  cvmean=cvmean+cvtemp
  cvmat[,i]=cvtemp/n
}
cvmean=cvmean/ndocv
cvmean=cvmean/n
plot(k_vec, cvmean, type="n", ylim=range(cvmat), xlab="k")

# plot each cv
for (i in 1:ndocv) lines(k_vec, cvmat[,i], col=i, lty=3)
# plot average
lines(k_vec, cvmean, type="b", col="black", lwd=3)
```


# Plot three cv curves on one figure

```
rgy = range(c(cv1,cv2,cv3))
plot(log(1/k_vec),cv1,type="l",col="red",ylim=rgy,lwd=2,
      cex.lab=0.8, xlab="log(1/k)", ylab="MSE")
lines(log(1/k_vec),cv2,col="blue",lwd=2)
lines(log(1/k_vec),cv3,col="green",lwd=2)
legend("topleft",legend=c("5-fold 1","5-fold 2","10 fold"),
      col=c("red","blue","green"),lwd=2,cex=0.8)
```



# Plot cross-validated MSE(k)

```
cv = (cv1+cv2+cv3)/3 #use average  
plot(k_vec, cv)
```




Get the best  $k$

```
k_best = k_vec[which.min(cv)]  
cat("the best k is: ", k_best, "\n")
```

```
## the best k is: 56
```

## Fit with the best $k$ and plot

```
near_best <- kknns(score~abv, train=dt1, test=dt1[,.(abv)],  
                  k=k_best, kernel='rectangular')  
plot(dt1$abv, dt1$score, cex.lab=1.2)  
lines(dt1$abv, near_best$fitted, col="red", lwd=2, cex.lab=2)
```



# Regression tree using cross-validation

```
big.tree=rpart(score~abv, data=dt1,  
               control=rpart.control(minsplit=5,  
                                     cp=0.0001,  
                                     xval=10))  
nbig = length(unique(big.tree$where))  
cat('size of big tree: ',nbig,'\n')
```

## size of big tree: 60

Find the best size of tree cp

```
cptable = printcp(big.tree)
```

##

## Regression tree:

## rpart(formula = score ~ abv, data = dt1, control = rpart.cont

## cp = 1e-04, xval = 10))

##

## Variables actually used in tree construction:


## [1] abv

##



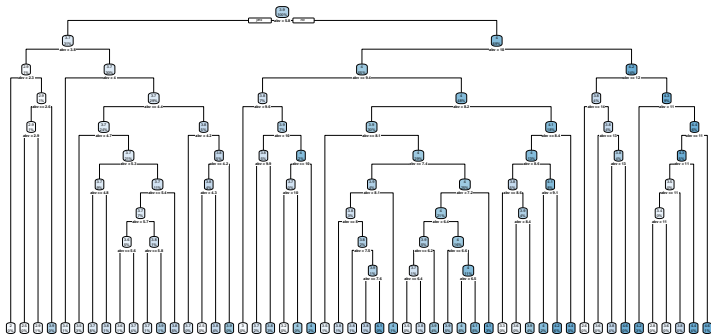
## Optimal cp parameter

```
bestcp = cptable[ which.min(big.tree$cptable[, "xerror"]), "CP" ]  
plotcp(big.tree) # plot results
```



# Prune down the optimal tree

```
par(mfrow=c(1,1))  
best.tree=prune(big.tree, cp=bestcp)  
rpart.plot(best.tree)
```




## Put everything together

```
dt1[, `:=`(lm_fit=yhat, knn_fit=near_best$fitted,  
           tree_fit=predict(best.tree))]
```

```
##           score  abv   lm_fit  knn_fit tree_fit  
##      1:    2.0  0.50 3.489640 3.000000 2.000000  
##      2:    2.0  2.20 3.594817 3.000000 2.000000  
##      3:    2.5  2.40 3.607191 3.000000 3.875000  
##      4:    5.0  2.40 3.607191 3.000000 3.875000  
##      5:    3.5  2.40 3.607191 3.000000 3.875000  
##      ---  
## 10475:    3.5 14.50 4.355802 3.785714 2.583333  
## 10476:    2.0 15.00 4.386737 3.785714 2.583333  
## 10477:    3.0 15.00 4.386737 3.785714 2.583333  
## 10478:    2.5 15.00 4.386737 3.785714 2.583333  
## 10479:    3.0 39.44 5.898808 3.767857 2.583333
```

# Plot the fit

```
plot(dt1$abv, dt1$score, pch=16, col='blue', cex=0.5)
lines(dt1$abv, dt1$lm_fit, col="red", lwd=2)
lines(dt1$abv, dt1$knn_fit, col="green", lwd=2)
lines(dt1$abv, dt1$tree_fit, col="yellow", lwd=2)
legend("bottomright", legend=c("y=score", "lm", "knn", "tree"),
      col=c("blue", "red", "green", "yellow"), lwd=2, cex=1.5)
```



# Pick the best performing model

- ▶ using test MSE

```
MSE <- dt1[-tr_idx,
           lapply(.SD,function(x) (x-score)^2),
           .SDcols=paste0(c("lm", "knn", "tree"),"_fit")]
MSE[,lapply(.SD, mean),
     .SDcols=paste0(c("lm", "knn", "tree"),"_fit")]
```

```
##          lm_fit  knn_fit  tree_fit
## 1: 0.4830103 0.4243779 0.4158278
```

- ▶ Tree does the best!

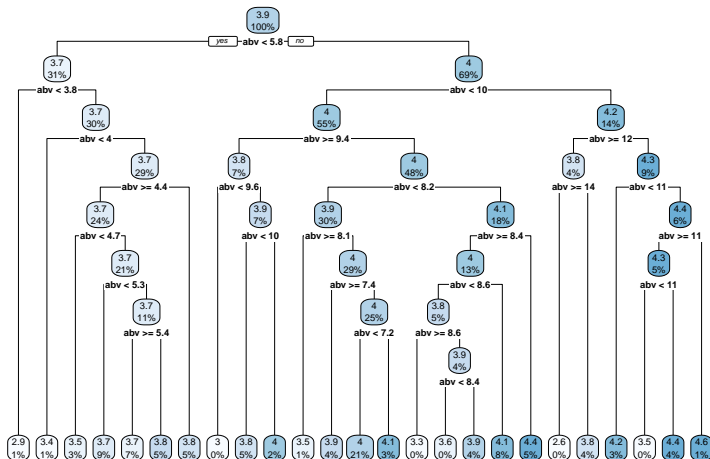
# Run regression trees with *rpart* package

Start with the big tree

```
temp = rpart(score~abv, data=dt1,  
             control=rpart.control(minsplit=5,  
                                   cp=0.001,  
                                   xval=0)  
             )
```

# Run regression trees with *rpart* package

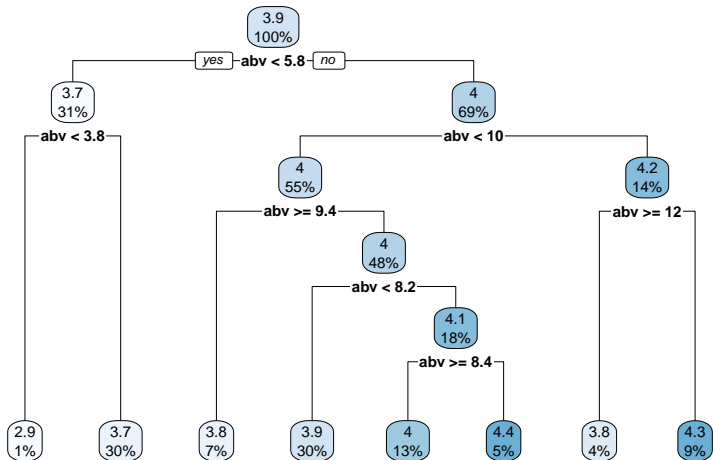
```
rpart.plot(temp)
```



# Run regression trees with *rpart* package

Prune the tree to have 8 leafs

```
beer.tree = prune(temp, cp=0.007)  
rpart.plot(beer.tree)
```






# Run regression trees with *rpart* package

Plot the data and fit


```
beer.fit = predict(beer.tree)
plot(dt1$abv, dt1$score, cex=.5, pch=16)
lines(dt1$abv, beer.fit, col="red", lwd=3) #step function fit
```



# Run regression trees with *rpart* package

Make prediction on two points

```
preddf = data.frame(abv=c(3,25))  
yhat = predict(beer.tree,preddf)  
plot(dt1$abv, dt1$score, cex=.5, pch=16) #plot data  
lines(dt1$abv,beer.fit,col="red",lwd=3)  
points(preddf$abv,yhat,col="blue",pch="*",cex=3)
```



## Run regression trees with *tree* package

```
temp2 <- tree(score~abv, data=dt1, mindev=0.0001)
cat("first big tree size: \n")
```

```
## first big tree size:
```

```
print(length(unique(temp2$where)))
```

```
## [1] 50
```

```
beer.tree2=prune.tree(temp2,best=8)
cat("pruned tree size: \n")
```

```
## pruned tree size:
```


```
print(length(unique(beer.tree2$where)))
```

```
## [1] 8
```

# Run regression trees with *tree* package

Plot the tree

```
plot(beer.tree2,type="uniform")  
text(beer.tree2,col="blue",label=c("yval"),cex=.8)
```



# k-NN and regression tree when x is multivariate

## ► Regression tree

```
temp2=rpart(score~., data=dt2,  
            control=rpart.control(minsplit=5,  
                                   cp=0.001,  
                                   xval=0))  
  
rpart.plot(temp2)  
  
# prune  
beer.tree2=prune(temp2, cp=0.008)  
rpart.plot(beer.tree2)
```

# k-NN and regression tree when x is multivariate

## ► Regression tree

```
# create a perspective plot
pv=seq(from=0.01,to=0.99,by=0.1)
x1q=quantile(dt2$abv, probs=pv)
x2q=quantile(dt2$age, probs=pv)
grids=expand.grid(x1q, x2q)
dt_pred <- data.table(age=grids[2], abv=grids[1])
colnames(dt_pred) <- c("age","abv")
beer.fit.plt=predict(beer.tree2,dt_pred)

persp(x1q, x2q, matrix(beer.fit.plt, ncol=length(x2q), byrow=T),
      theta=150, xlab='age', ylab='abv', zlab='score',
      zlim=c(min(dt2$score), 1.1*max(dt2$score)))
```

# k-NN and regression tree when x is multivariate

## ► k-NN

```
# scale input variables
```

```
scf = function(x) {return((x-min(x))/(max(x)-min(x)))}  
dt2[, `:=` (abv=scf(abv), age=scf(age))]
```

```
##           score          abv          age  
##      1:    2.0 0.00000000 0.08702472  
##      2:    2.0 0.04365691 0.20658398  
##      3:    2.5 0.04879301 0.06580783  
##      4:    5.0 0.04879301 0.12898571  
##      5:    3.5 0.04879301 0.20658398  
##      ---  
## 10475:    3.5 0.35952748 0.03217994  
## 10476:    2.0 0.37236775 0.11320478  
## 10477:    3.0 0.37236775 0.13974191  
## 10478:    2.5 0.37236775 0.18066741  
## 10479:    3.0 1.00000000 0.16997031
```

# k-NN and regression tree when x is multivariate

## ► k-NN

```
beer.knn <- kknnscore~., train=dt2, test=dt2[,.(abv, age)],  
           k=5, kernel='rectangular')  
dt1[,knn_fit:=beer.knn$fitted.values]
```

##		score	abv	lm_fit	knn_fit	tree_fit
##	1:	2.0	0.50	3.489640	2.8	2.000000
##	2:	2.0	2.20	3.594817	2.4	2.000000
##	3:	2.5	2.40	3.607191	2.2	3.875000
##	4:	5.0	2.40	3.607191	3.0	3.875000
##	5:	3.5	2.40	3.607191	2.4	3.875000
##	---					
##	10475:	3.5	14.50	4.355802	3.4	2.583333
##	10476:	2.0	15.00	4.386737	2.5	2.583333
##	10477:	3.0	15.00	4.386737	2.4	2.583333
##	10478:	2.5	15.00	4.386737	2.8	2.583333
##	10479:	3.0	39.44	5.898808	2.8	2.583333



# Classification using k-NN

- Use fgl dataset from R library

Load the libraries and dataset

```
library("MASS")  
library("kkn")  
data(fgl)
```

View dataset codebook

```
help(fgl)
```

Quick look at the dataset

```
head(fgl,n=3)
```

##		RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	type
## 1		3.01	13.64	4.49	1.10	71.78	0.06	8.75	0	0	WinF
## 2		-0.39	13.89	3.60	1.36	72.73	0.48	7.83	0	0	WinF
## 3		-1.82	13.53	3.55	1.54	72.99	0.39	7.78	0	0	WinF

# Classification using k-NN

Code the 7 types into 3 main categories

```
n=nrow(fgl)
y = rep(3,n)
y[fgl$type=="WinF"]=1
y[fgl$type=="WinNF"]=2
y = as.factor(y)
levels(y) = c("WinF","WinNF","Other")
print(table(y,fgl$type))
```

```
##
## y          WinF WinNF Veh Con Tabl Head
##   WinF      70     0   0   0    0    0
##   WinNF     0    76   0   0    0    0
##   Other     0     0  17  13    9   29
```

```
x = cbind(fgl$RI,fgl$Na,fgl$A1)
```

# Classification using k-NN

Scale the input variables, construct ready-to-use dataset

```
scf = function(x) {return((x-min(x))/(max(x)-min(x)))}  
x = apply(x,2,scf)      # scale all x  
  
colnames(x) = c("RI", "Na", "Al")  
ddf = data.frame(type=y,x)  
names(ddf) =c("type", "RI", "Na", "Al")
```

Check the dataset

```
tail(ddf, n=3)
```

##	type	RI	Na	Al
## 212	Other	0.4170325	0.5458647	0.5389408
## 213	Other	0.2352941	0.5488722	0.5140187
## 214	Other	0.2616330	0.5263158	0.5576324

# Classification using k-NN

## ► Run k-NN

```
near = kknnc(type~.,ddf,ddf,k=10,kernel = "rectangular")
print(table(near$fitted,ddf$type))
```

```
##
##           WinF WinNF Other
##   WinF      58   13   14
##   WinNF     11   57   12
##   Other      1    6   42
```

Predicted probabilities

```
fitdf = data.frame(type=ddf$type,near$prob)
names(fitdf)[2:4] = c("ProbWinF","ProbWinNF","ProbOther")
head(fitdf,n=3)
```

```
##   type ProbWinF ProbWinNF ProbOther
## 1 WinF      0.6      0.3      0.1
## 2 WinF      0.4      0.4      0.2
## 3 WinF      0.1      0.9      0.0
```

# Classification using k-NN

```
par(mfrow=c(1,3))  
plot(ProbWinF~type,fitdf,col=c(grey(.5),2:3),cex.lab=1.4)  
plot(ProbWinNF~type,fitdf,col=c(grey(.5),2:3),cex.lab=1.4)  
plot(ProbOther~type,fitdf,col=c(grey(.5),2:3),cex.lab=1.4)
```

