

Recommender Systems

Mladen Kolar

Tools so far

We already know many things that we can use to recommend new content.

- ▶ predict rating of a restaurant
- ▶ predict if a user is going to respond to a marketing campaign
- ▶ predict if user is going to like a music

Why do we need recommendations?

Help users discover new content

Because you watched Justice League



Because you watched Firefly



Why do we need recommendations?

Help users find what they have been already looking for.

Prominently place content as a part of the user interface.



The image shows a screenshot of an Amazon product page for the movie "Harry Potter and the Sorcerer's Stone". The main visual is a large poster of the first Harry Potter film. At the top, the title "Harry Potter and the Sorcerer's Stone" is displayed along with the year "2001", rating "PG-13", and closed captioning "CC". Below the poster, there is a rating of "★★★★★ 2,829" and an IMDB score of "7.5/10". A "Watch Trailer" button is visible. The plot summary describes Harry as an orphaned boy who discovers he has magical powers on his eleventh birthday. Below the summary, it lists "Starring: Richard Harris, Maggie Smith" and "Runtime: 2 hours, 33 minutes". It also mentions that the movie is available to watch on "supported devices". On the right side of the page, there are buttons for "Rent or Buy", "Rent HD \$3.99", "Buy HD \$13.99", and "Redeem a gift card". At the bottom, there are "Add to Wish List" and "Add to Cart" buttons. Social sharing icons for email, Facebook, and Twitter are located at the bottom left. A small note at the bottom right states: "By placing your order, you agree to our [Terms of Use](#). Sold by Amazon Digital Services, Inc."

Customers Who Watched This Item Also Watched



Why do we need recommendations?

Help users find complementary products



Levi's Men's 501 Original Shrink To Fit Jean

★★★★☆ • 1,353 customer reviews | 6 answered questions

Price: \$17.33 - \$59.99 & FREE Returns on some sizes and colors. [Details](#)

Sale: Lower price available on select options

Fit: As expected (71%)

Size:

Select [Size Chart](#)

Color: Rigid STF



- 100% Cotton
- Imported
- Machine Wash
- Straight-leg jeans in shrink-to-fit denim that fades with repeat washings
- Five-pocket styling
- Button Fly
- 17.25-leg opening
- Actual coloration may vary from garment to garment due to specific wash processor

Customers Who Bought This Item Also Bought



Levi's Men's 40MM
Reversible Belt With
Gunmetal Buckle
★★★★☆ 555
\$13.99 - \$19.99



Carhartt Men's Twill Long
Sleeve Work Shirt Button
Front
★★★★☆ 132
\$24.99 - \$25.99



Wrangler Men's Sport
Western Snap Long Sleeve
Shirt
★★★★☆ 147
\$19.95 - \$26.50



Burton Men's Locked
Long-Sleeve Woven Shirt
★★★★☆ 76
\$15.01 - \$42.00



Levi's Men's 514 Straight-
Fit Corduroy Pant
★★★★☆ 30
\$42.99



U.S. Polo Assn. Men's Slim
Fit Solid Pique Polo Shirt
★★★★☆ 71
\$11.82 - \$30.99



Clarks Originals Men's
Desert Boot
★★★★☆ 2,643
#1 Best Seller In Men's
Chukka Boots
\$56.88 - \$199.00

Why do we need recommendations?

Help users find substitute products



Levi's Men's 501 Original Shrink To Fit Jean

★☆☆☆☆ 1,353 customer reviews | 6 answered questions

Price: \$17.33 - \$59.99 & FREE Returns on some sizes and colors. [Details](#)

Sale: Lower price available on select options

Fit: As expected (71%)

Size:

Select Size Chart

Color: Rigid STF



- 100% Cotton
- Imported
- Machine Wash
- Straight-leg jean in shrink-to-fit denim that fades with repeat washings
- Five-pocket styling
- Button Fly
- 17.25-inch opening
- Actual coloration may vary from garment to garment due to specific wash processor

Customers Who Viewed This Item Also Viewed



Levi's Men's Jeans 501
Original Fit
★☆☆☆☆ 4,655
\$34.99 - \$64.99



Levi's Men's 505 Regular
Fit Jean
★☆☆☆☆ 5,479
#1 Best Seller In Men's
Jeans
\$39.49 - \$68.00



Levi's Men's 514 Straight
Jean
★☆☆☆☆ 3,089
\$19.15 - \$56.00



Levi's Men's Big-Tall 501
Shrink To Fit Jean
★☆☆☆☆ 253
\$39.99 - \$59.99



Levi's Men's 501 Colored
Rigid Shrink-to-Fit Jean
(Clearance), Light Gray
★☆☆☆☆ 63
\$16.70 - \$49.99



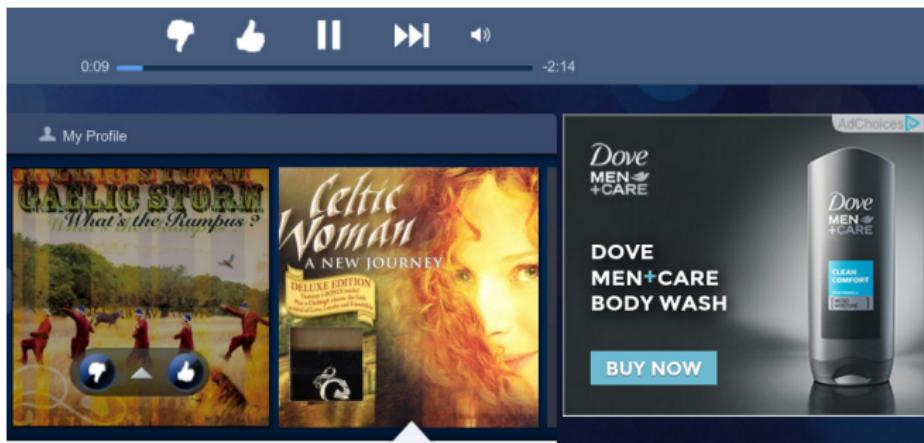
Levi's Men's 501 Colored
Rigid Shrink-to-Fit Jean
(Clearance), Blue Green
Rigid
★☆☆☆☆ 80
\$16.83 - \$49.99



Why do we need recommendations?

Personalize user experience based on the feedback

Recommend products based on our interests



Why do we need recommendations?

Who to follow?

 Joe Blitzstein @stat110 Statistics professor at Harvard; statistician and data scientist; probability and paradoxes; Bayesian frequentist reconciliation; chess. Followed by Bert Huang, Sherri Rose and Hal Daumé III.	 
 Sean J. Taylor @seanjtaylor Social Scientist. Hacker. Facebook Data Science Team. Keywords: Experiments, Causal Inference, Statistics, Machine Learning, Economics. Followed by José Pablo González, Sherri Rose and John D. Cook.	 
 Daniel Roy @roydanroy Asst Professor Of Stats at UofT working also in theoretical computer science, machine learning, and probability. I like randomness. Followed by Arthur Gretton, Dino Sejdinovic and Neil Lawrence.	 
 Jason Baldridge @jasonbaldridge Co-founder of @PeoplePattern and Associate Professor of Computational Linguistics, UT Austin. Followed by José Pablo González, Jonathan Clark and petricek.	 
 Nando de Freitas @NandoDF Oxford University Prof & Google DeepMind Scientist	 
 Ryan Adams @ryan_p_adams Computer Science Professor, Machine Learning Researcher @Twitter, Entrepreneur, Podcaster, Dad, Sports Fan Followed by Bert Huang, Hal Daumé III and Tomasz Malisiewicz.	 

Why do we need recommendations?

Help us find things we like

Results for 'the guild'

The Guild: Seasons 1 & 2
2007 NR 2 discs / 2 episodes

Actress and former "World of Warcraft" addict Felicia Day created and stars in this Web series about a quirky group of gamers who take part in all sorts of heroics online but are far less skilled at navigating real life.

Starring: Felicia Day, Jeff Lewis
Creator: Felicia Day
Genre: TV Sitcoms
Format: DVD and streaming (HD available)

Play + All

★★★☆☆ Our best guess for Mike

A screenshot of a digital media platform's search results page. The search term 'the guild' is entered. The top result is 'The Guild: Seasons 1 & 2' from 2007, rated NR, containing 2 discs with 2 episodes. A brief description highlights the show's focus on gamers. Below the main title, it lists the starring actors (Felicia Day, Jeff Lewis), creator (Felicia Day), genre (TV Sitcoms), and format (DVD and streaming, with HD available). At the bottom of the card, there are two buttons: 'Play' and '+ All'. To the right of the card, there is a rating section showing three red stars and one yellow star, followed by the text 'Our best guess for Mike'. On the left side of the card, there is a small thumbnail image of the show's cast.

Why do we need recommendations?

- ▶ Discover new content
- ▶ Help us find what we are looking for
- ▶ Personalize content based on our feedback/interests
- ▶ Help us find things we like
- ▶ ...

All of the above problems are slightly different from each other.

However, at the base of all of them is the need to model preferences, opinions and behaviour of users.

Popularity

What are people viewing now

Limitations

- ▶ no context (what is my intention now)
- ▶ no personalization

Trends · Change

#NationalStressAwarenessDay
77.7K Tweets about this trend

#YouHadMeAt
26.5K Tweets about this trend

#FoxLake
Just started trending

#BAYvAFC
42.7K Tweets about this trend

San Diego
40.3K Tweets about this trend

MOST EMAILED	MOST VIEWED
1. Missing Italian Marathoner Found on New York Subway, Still in His Running Gear	
2. Beyond the Honeycrisp Apple	
3. <small>ESSAY</small> On the Longest Hiking Trails, a Woman Finds Equal Footing	
4. <small>THE UPSHOT</small> Stressed, Tired, Rushed: A Portrait of the Modern Family	
5. <small>ABOUT NEW YORK</small> Sideline to a Spy Saga: How a Brooklyn Newsboy's Nickel Would Turn Into a Fortune	
6. DAVID BROOKS The Evolution of Simplicity	
7. FRANK BRUNI The Catholic Church's Sins Are Ours	
8. Robin Williams's Widow Points to Dementia as a Suicide Cause	
9. OP-ED THOMAS B. EDGALL Why Are Asian-Americans Such Loyal Democrats?	
10. Review: Resettling the Meaning of Home in Brooklyn, 'With Saoirse Ronan'	

[Go to Complete List »](#)

Predicting rating using existing models

Build a model using features of a user and an item



David L. Alberta

Grand Rapids, MI

Reviewer ranking: #1,450,021

89% helpful

votes received on reviews
(163 of 184)

[Send an Email](#)

PUBLIC ACTIVITY

[Reviews \(15\)](#)

OTHER

[Give Profile Feedback to Amazon >](#)

Reviewed



TrollHunter

120 of 128 people found this review helpful

★★★★★ Quirky, Intelligent and just plain fun to watch

June 26, 2011

This film demonstrates how mockumentaries should be made. The film never takes itself too serious and is just plain fun. The trolls look like doofy monsters from the Grimms Fairy-Tales and are not like the monsters you see in animated kid flicks. A Troll Hunter explains, fairy-tales and real life are not the same. Trolls are just dumb animals and as animals they act on their instinctive behavior. And why are the trolls in the movie being hunted down? Well, you will see if you watch the movie.

[...Read more >](#)

Product Details

Genres	Fantasy, Drama, International, Comedy, Horror
Director	Andre Øvredal
Starring	Otto Jespersen, Glenn Erland Tostnerud
Supporting actors	Johanna Merck, Tomas Alf Larsen, Urmila Berg-Domaas, Hans Morten Hansen, Robert Stoltzenberg, Knut Nærum, Eirik Bech, Inge Erik Herjesand, Tom Jørgensen, Benedicte Auber Ringnes, Magne Skjæveland, Torunn Ledemel Stokkeland, Finn Norvald Øvredal, Kaja Halden Arnestad, Robin De Lano, Jens Stoltzenberg
Studio	Magnolia
MPAA rating	PG-13 (Parental Guidance Suggested)
Captions and subtitles	English <small>Details</small> <small>▼</small>
Rental rights	48 hour viewing period. Details <small>▼</small>
Purchase rights	Stream instantly and download to 2 locations Details <small>▼</small>
Format	Amazon Video (streaming online video and digital download)

Predicting rating using existing models

The Music Genome Project

<https://www.pandora.com/about/mgp>

... Each song in the Music Genome Project is analyzed using up to 450 distinct musical characteristics by a trained music analyst. These attributes capture not only the musical identity of a song, but also the many significant qualities that are relevant to understanding the musical preferences of listeners. The typical music analyst working on the Music Genome Project has a four-year degree in music theory, composition or performance, has passed through a selective screening process and has completed intensive training in the Music Genome's rigorous and precise methodology. ...

Predicting rating using existing models



[Lady Gaga](#)

[Just Dance \(Remix Single\)](#)

PANDORA®

Just Dance (Redone Remix F. Kardinal Offishall)

[Play Sample](#)



Features Of This Track

electronica roots
trip hop roots
r&b influences
funk influences
beats made for dancing
unsyncopated ensemble rhythms
straight drum beats
a female vocal
clear pronunciation
a rhythmic intro
use of modal harmonies
the use of chordal patterning
melodic part writing
use of strings
subtle use of arpeggiated synths
affected synths

Obtaining features may be expensive

Predicting rating using existing models

Build a model using features of a user and an item

- ▶ linear model
- ▶ decision tree
- ▶ boosting model
- ▶ ...

$$\text{rating} = f(\text{user features}, \text{item features})$$

Predicting rating using existing models

We could predict how many stars a user would give.

We could predict whether a user would give a high rating if they purchase an item.

We could predict whether a user will like an item.

Based on the predicted rating, the system would recommend

- ▶ diverse set of items
- ▶ new/unseen content
- ▶ something that the user likes or searches for

Predicting rating using existing models

We train a model based on user's past feedback and various features that characterize users and items.

Recommender system provides a rating and we provide recommendations based on this.

Similar to the example in marketing where customers are targeted based on \hat{p} .

Predicting rating using existing models

Approach does not suffer from a cold-start problem

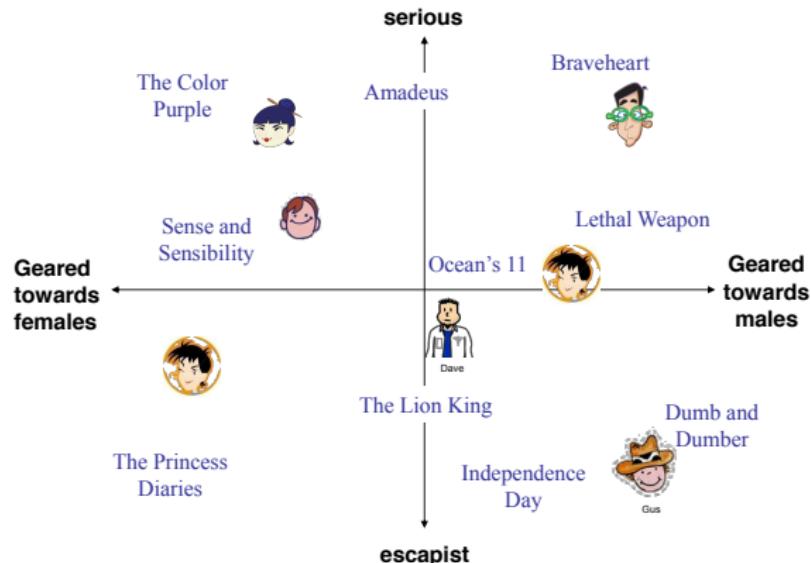
- ▶ Rate new movie from features of other movies user liked

Limitations

- ▶ sometimes, we do not have access to features
- ▶ often does not perform as well as collaborative filtering methods

Learning relationships

Recommender systems can uncover/model relationships between users and items they are evaluating, using historical ratings.



Collaborative filtering

Data we may have (example from Netflix challenge)

Training data

user	movie	date	score
1	21	5/7/02	1
1	213	8/2/04	5
2	345	3/6/01	4
2	123	5/1/05	4
2	768	7/15/02	3
3	76	1/22/01	5
4	45	8/3/00	4
5	568	9/10/05	1
5	342	3/5/03	2
5	234	12/28/00	2
6	76	8/11/02	5
6	56	6/15/03	4

Test data

user	movie	date	score
1	62	1/6/05	?
1	96	9/13/04	?
2	7	8/18/05	?
2	3	11/22/05	?
3	47	6/13/02	?
3	15	8/12/01	?
4	41	9/1/00	?
4	28	8/27/05	?
5	93	4/4/05	?
5	74	7/16/03	?
6	69	2/14/04	?
6	83	10/3/03	?

Collaborative filtering

Data commonly represented as a rating matrix.

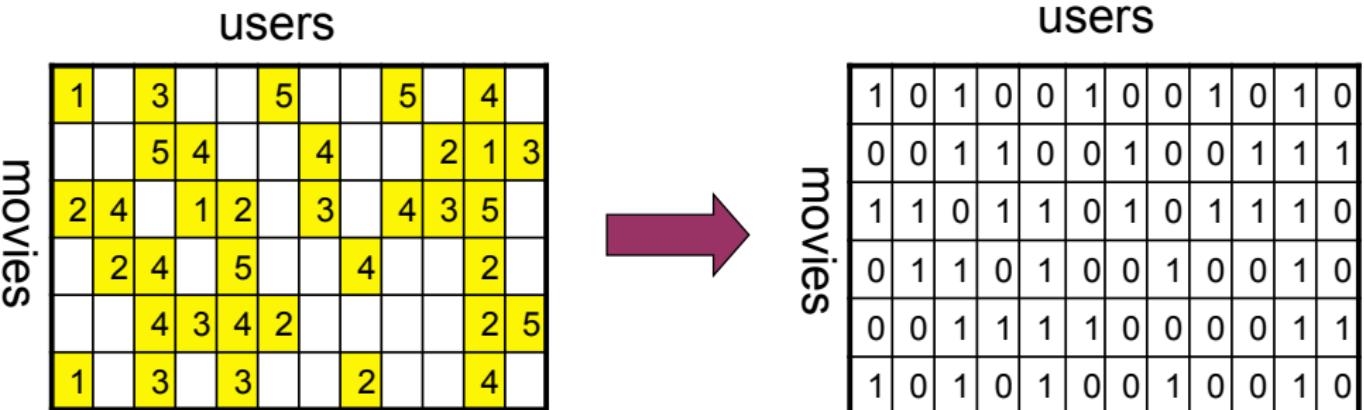
		users											
		1	2	3	4	5	6	7	8	9	10	11	12
items	1	1		3			5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	


- unknown rating


- rating between 1 to 5

Collaborative filtering

- ▶ a user viewed/bought an item (1)
- ▶ a user did not view/buy an item (0)



Collaborative filtering

- ▶ 1 user viewed/bought an item and liked it
 - ▶ -1 user viewed/bought an item and did not like it
 - ▶ 0 user did not view/buy an item

users

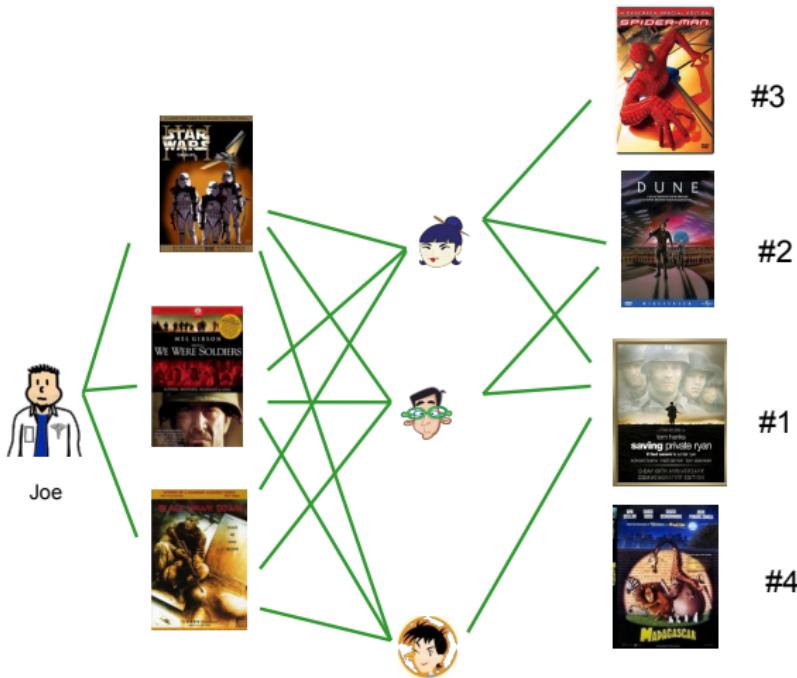
1		3		5		5		4
		5	4		4		2	1
2	4		1	2	3	4	3	5
	2	4		5		4		2
	4	3	4	2			2	5
1		3	3		2		4	



users

-1	0	-1	0	0	1	0	0	1	0	1	0
0	0	1	1	0	0	1	0	0	-1	-1	-1
-1	1	0	-1	-1	0	-1	0	1	-1	1	0
0	-1	1	0	1	0	0	1	0	0	-1	0
0	0	1	-1	1	-1	0	0	0	0	-1	1
-1	0	-1	0	-1	0	0	-1	0	0	1	0

Basic idea



Basic idea

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
items	1	1		3		??	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	
	7												

 - unknown rating  - rating between 1 to 5

Basic idea

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
items	1	1		3		??	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	

 - unknown rating  - rating between 1 to 5

similarity

$$s_{13} = 0.2$$

$$s_{16} = 0.3$$

Predict rating using weighted average

$$\frac{0.2 \cdot 2 + 0.3 \cdot 3}{0.2 + 0.3} = 2.6$$

Basic idea

(user, user) similarity to recommend items

(item, item) similarity to recommend new items that were also liked by the same users

The oldest known collaborative filtering method.

See Amazon recommendation system at scale:

<http://www.cs.umd.edu/~samir/498/Amazon-Recommendations.pdf>

Similarities

How do we measure (user, user) similarity or (item, item) similarity?

- ▶ Euclidean distance
- ▶ Jaccard similarity
- ▶ Cosine similarity
- ▶ Pearson correlation

Notation:

$$\text{rating}(\text{user}, \text{item}) = r(u, i) = r_{ui}$$

I_u = set of items purchased by user u

U_i = set of users who purchased by item i

Euclidean distance

Distance between item i_1 and item i_2 is

$$\text{dist}^2(i_1, i_2) = \sum_u (r_{u,i_1} - r_{u,i_2})^2$$

If each rating is 0 or 1 (user bought item or did not), then the distance becomes

$$\begin{aligned}\text{dist}^2(i_1, i_2) &= \#\{\text{users that bought } i_1, \text{ but not } i_2\} \\ &\quad + \#\{\text{users that bought } i_2, \text{ but not } i_1\}\end{aligned}$$

Euclidean distance

Example:

$$U_1 = \{1, 4, 8, 9, 11, 23, 25, 34\}$$

$$U_2 = \{1, 4, 6, 8, 9, 11, 23, 25, 34, 35, 38\}$$

$$U_3 = \{4\}$$

$$U_4 = \{5\}$$

What is the distance between items 1 and 2?

What is the distance between items 3 and 4?

Euclidean distance

What is the distance between items 1 and 2? $\rightarrow 3$

What is the distance between items 3 and 4? $\rightarrow 2$

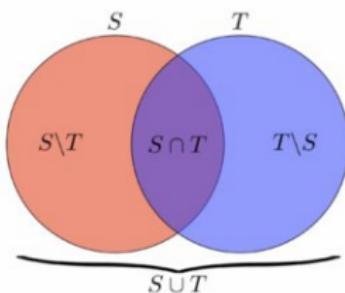
Problem:

- ▶ Favors small sets, even if they have few elements in common.

Jaccard similarity

Measures similarity between sets

$$\begin{aligned}\text{Jaccard}(U_i, U_j) &= \frac{|U_i \cap U_j|}{|U_i \cup U_j|} \\ &= \frac{\text{bought i and j}}{\text{bought i or j}}\end{aligned}$$



Key idea: normalize by popularity

Jaccard similarity

- ▶ Maximum of 1 if two items were purchased by the same set of users or if the two users purchased exactly the same set of items.
- ▶ Minimum of 0 if the two items were purchased by completely disjoint sets of users or if the two users purchased completely disjoint sets of items.

Jaccard similarity in action

How does amazon generate their recommendations?

<http://www.cs.umd.edu/~samir/498/Amazon-Recommendations.pdf>

A user is looking at



U_i is the set of users who viewed this product.

Rank products according to $\frac{|U_i \cap U_j|}{|U_i \cup U_j|}$



.86



.84



.82



.79



...



...



...



Cosine similarity

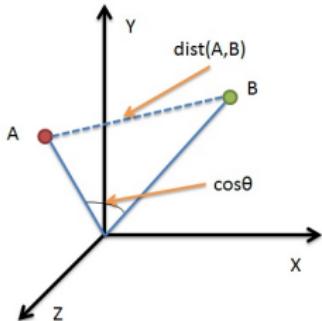
Jaccard similarity works only on 0/1 data

Cosine similarity works on arbitrary data.

users												
movies	-1	0	-1	0	0	1	0	0	1	0	1	0
	0	0	1	1	0	0	1	0	0	-1	-1	-1
	-1	1	0	-1	-1	0	-1	0	1	-1	1	0
	0	-1	1	0	1	0	0	1	0	0	-1	0
	0	0	1	-1	1	-1	0	0	0	0	-1	1
	-1	0	-1	0	-1	0	0	-1	0	0	1	0

- ▶ 1 user viewed/bought an item and liked it
- ▶ -1 user viewed/bought an item and did not like it
- ▶ 0 user did not view/buy an item

Cosine similarity

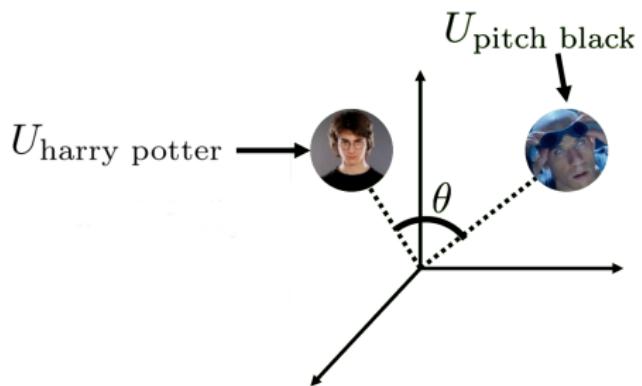


$$\text{similarity}(A, B) = \cos(\theta) = \frac{A \cdot B}{\|A\| \cdot \|B\|}$$

- ▶ $\cos(\theta) = 1$ ($\theta = 0$) A and B point in the same direction
- ▶ $\cos(\theta) = -1$ ($\theta = 180$) A and B point in the opposite direction
- ▶ $\cos(\theta) = 0$ ($\theta = 90$) A and B are orthogonal

Cosine similarity

Each item is represented by a vector of users' ratings



$$U_{\text{harry potter}} = (0, 1, 1) \quad U_{\text{pitch black}} = (1, 1, 0)$$

$$\text{similarity} = \frac{(0, 1, 1) \cdot (1, 1, 0)}{\sqrt{2} \cdot \sqrt{2}} = \frac{1}{2}$$

Pearson correlation

Very similar to cosine similarity

Cosine similarity would fail if naively applied to ratings

Example:

$$\begin{aligned} R_1 &= (1, 1, 1) \\ R_2 &= (5, 5, 5) \end{aligned} \longrightarrow \text{similarity} = 1$$

Pearson correlation will solve this by removing average from the rating vector

Pearson correlation

User ratings for item i:

1	?	?	5	5	3	?	?	?	4	2	?	?	?	?	4	?	5	4	1	?
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

User ratings for item j:

?	?	4	2	5	?	?	1	2	5	?	?	2	?	?	3	?	?	?	5	4
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Pearson correlation computed over shared support

$$s_{ij} = \frac{\sum_{u \in U_i \cap U_j} (r_{ui} - \bar{r}_i)(r_{uj} - \bar{r}_j)}{\sqrt{\sum_{u \in U_i \cap U_j} (r_{ui} - \bar{r}_i)^2 \cdot \sum_{u \in U_i \cap U_j} (r_{uj} - \bar{r}_j)^2}}$$

Pearson correlation vs cosine similarity

Pearson correlation

$$s_{ij} = \frac{\sum_{u \in U_i \cap U_j} (r_{ui} - \bar{r}_i)(r_{uj} - \bar{r}_j)}{\sqrt{\sum_{u \in U_i \cap U_j} (r_{ui} - \bar{r}_i)^2 \cdot \sum_{u \in U_i \cap U_j} (r_{uj} - \bar{r}_j)^2}}$$

Cosine similarity

$$s_{ij} = \frac{\sum_{u \in U_i \cap U_j} r_{ui} \cdot r_{uj}}{\sqrt{\sum_{u \in U_i \cap U_j} r_{ui}^2 \cdot \sum_{u \in U_i \cap U_j} r_{uj}^2}}$$

How to use similarity to recommend?

Given similarity measure between items s_{ij}

Find set $s_k(i, u)$ of k -nearest neighbors to movie i that were rated by user u

Estimate rating using weighted average over the set of neighbors

$$\hat{r}_{ui} = \frac{\sum_{j \in s_k(i, u)} s_{ij} r_{uj}}{\sum_{j \in s_k(i, u)} s_{ij}}$$

Normalization and bias problem

Problems:

- ▶ Some items are significantly higher rated
- ▶ Some users rate substantially lower
- ▶ Ratings change over time

Bias correction is crucial for collaborative filtering approaches

- ▶ global bias μ
- ▶ offset per user b_u
- ▶ offset per movie b_i
- ▶ time effects (ignore for now)

Baseline rating for (user, movie) is $b(u, i) = \mu + b_u + b_i$

Normalization and bias problem



Mean rating of all movies $\mu = 3.7$

Troll hunter is 0.7 above mean ($b_i = 0.7$)

User rates 0.2 below mean ($b_u = -0.2$)

For this (user, movie) baseline rating is 4.2 stars

Estimating biases

$$\min_b \sum_{(u,i):r(u,i)\neq?} (r(u,i) - \mu - b_u - b_i)^2 + \lambda(\sum_u b_u^2 + \sum_i b_i^2)$$

This is a linear model. Why?

How to recommend with biases?

Similar to the approach earlier

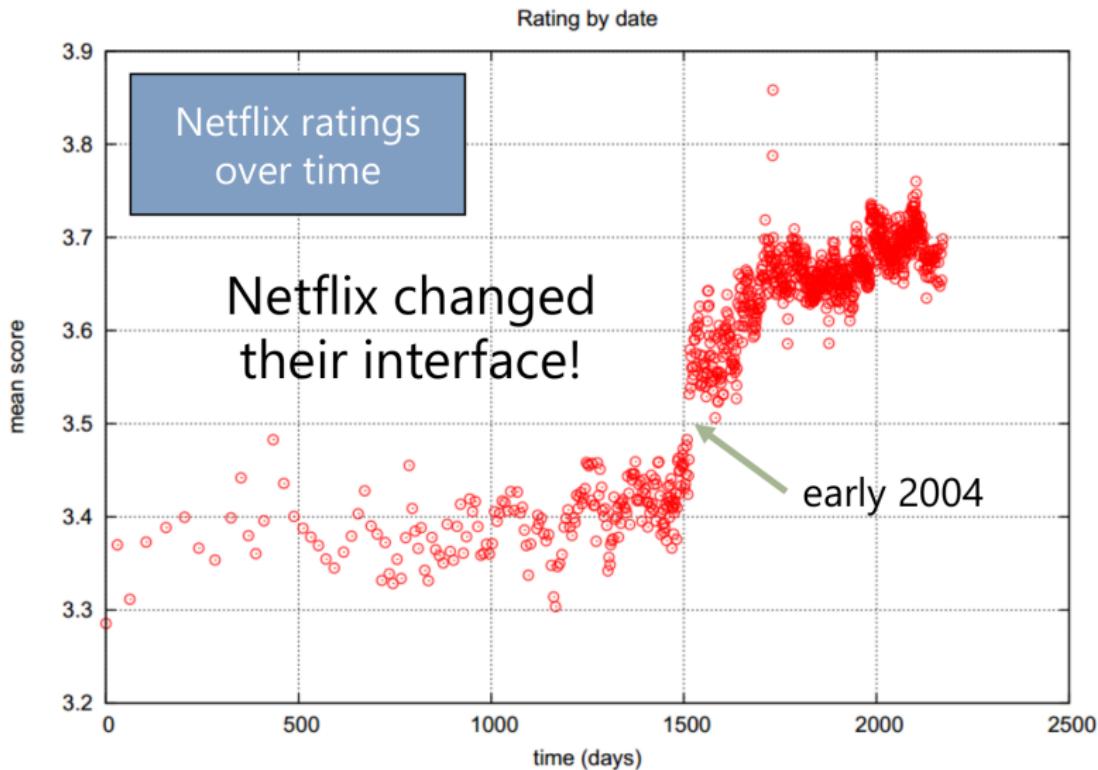
Given similarity measure between items s_{ij}

Find set $s_k(i, u)$ of k -nearest neighbors to movie i that were rated by user u

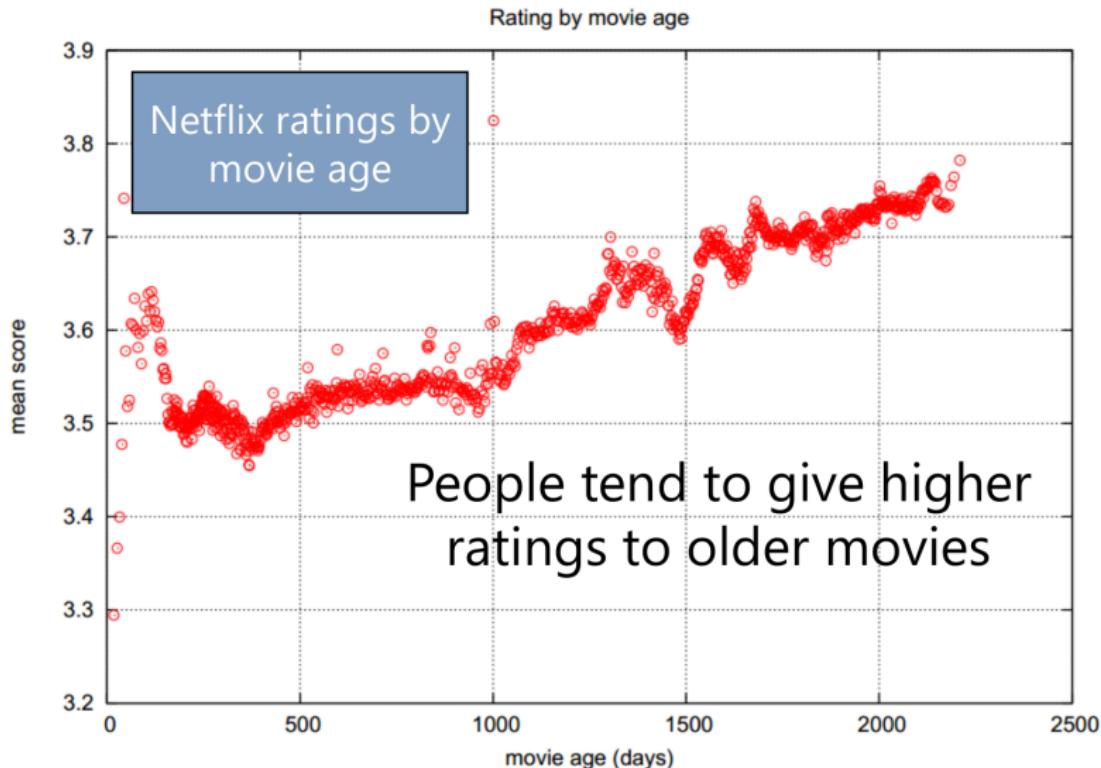
Estimate rating using weighted average over the set of neighbors

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in s_k(i, u)} s_{ij} (r_{uj} - b_{uj})}{\sum_{j \in s_k(i, u)} s_{ij}}$$

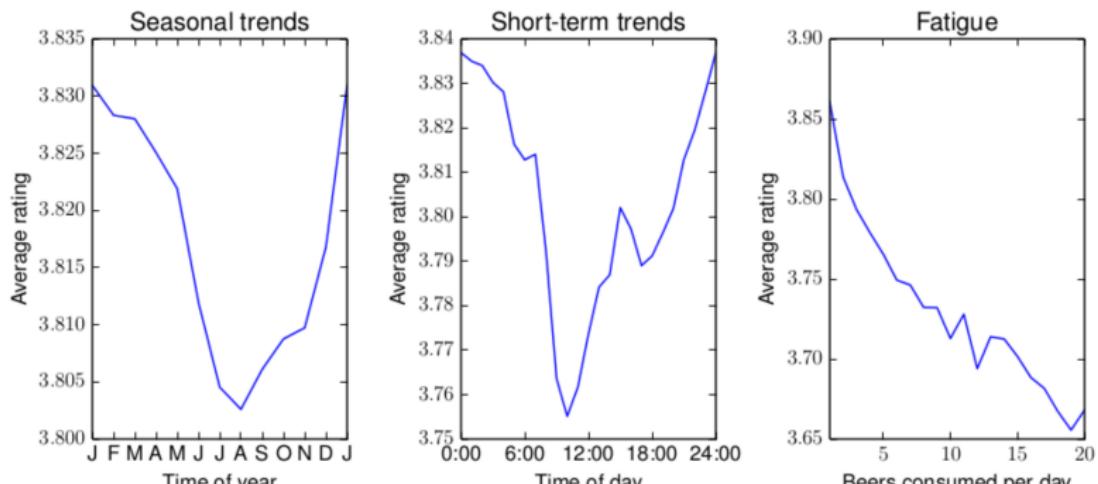
Temporal effects



Temporal effects



Temporal effects



A few temporal effects from beer reviews

Recommendations based on similarities

- ▶ intuitive
- ▶ there is no training
- ▶ easy to explain to a user
- ▶ widely used in practice
- ▶ surprisingly, collaborative filtering is extremely useful even though we have not looked at any features
- ▶ accuracy and scalability questionable
- ▶ cold start problem
- ▶ will not necessarily encourage diverse results

Cold start problem

What happens with new users where we have no ratings yet?

- ▶ Recommend popular items.
- ▶ Have some start-up questions (for example, “tell me 10 movies you love”).

What do we do with new items?

- ▶ Content-based filtering techniques (that is, use features).
- ▶ Pay a focus group to rate them.

Modelling approach to recommendations

So far we've looked at approaches that try to define some definition of user/user and item/item similarity.

Recommendation consists of

1. Finding an item i that a user likes (gives a high rating)
2. Recommending items that are similar to it (that is, items j with a similar rating profile to i)

Hinges on finding a good measure of similarity

Modelling approach to recommendations

What we want to do next is to model the rating.

$$r_{ui} = f(\text{user}, \text{item}) + \text{noise}$$

Recommendation consists of identifying items with largest rating

$$\text{recommendation}(u) = \arg \max_{i \in \text{unseen items}} f(u, i)$$

Netflix yardstick

Netflix prize

- ▶ In 2006, Netflix created a dataset of 100,000,000 movie ratings
- ▶ Data looked like: (userID, itemID, time, rating)
- ▶ Whoever first manages to reduce the (R)MSE by 10% versus Netflix's solution wins \$1,000,000
- ▶ Data were de-anonymized — lawsuit against Netflix

Root mean squared error for predicting ratings

$$\text{RMSE}(f) = \sqrt{\frac{1}{N} \sum_{u,i} (f(u, i) - r_{ui})^2}$$

Predicted rating $\hat{r}_{ui} = f(u, i)$

Netflix yardstick

A lot of research on minimizing the mean squared error.

Not clear that improving this metric will lead to better user experience.

When building a model we focus on minimizing root mean squared error.

Latent factor model

Suppose we had K features of movies and users

Describe movie i with features q_i

- ▶ How much is it action, romance, drama, ...

$$q_i = (0.9, 0.2, 0.5, \dots)$$

Describe user u with features p_u

- ▶ How much she likes action, romance, drama, ...

$$p_u = (0.01, 0, 0.9, \dots)$$

$f(u, i)$ is the product of the two vectors

$$f(u, i) = 0.9 \cdot 0.01 + 0.2 \cdot 0 + 0.5 \cdot 0.9 + \dots$$

Discovering features via matrix factorization

users

1		3		5		5	4	
		5	4		4		2	1
	2	4	1	2	3	4	3	5
		2	4	5		4		2
		4	3	4	2			2
1		3	3		2		4	

items

~

users

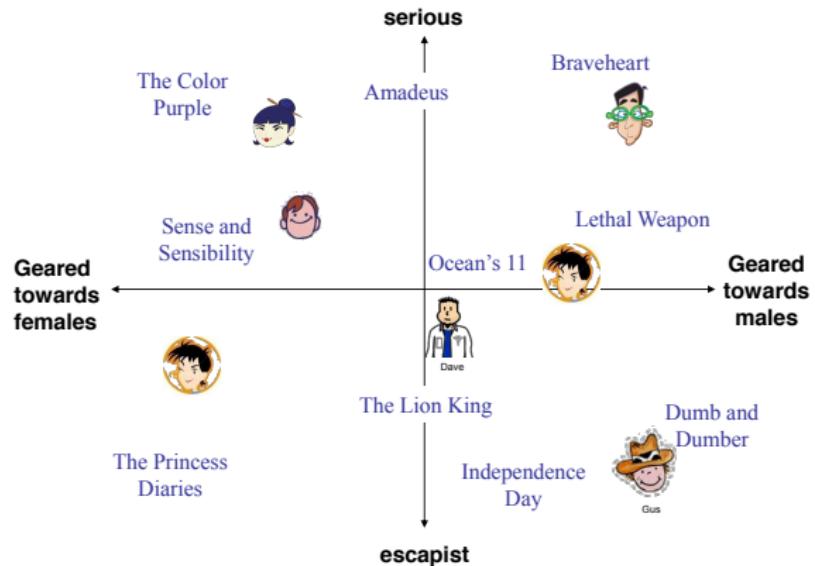
1.1	-2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

items

~

.1	-.4	.2
-.5	.6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-2
-1	.7	.3

Discovering features via matrix factorization



Discovering features via matrix factorization

users

1		3		5		5	4	
		5		?	4		2	1 3
2	4		1 2		3	4	3	5
	2	4		5		4		2
		4	3	4	2			2 5
1		3		3		2		4

items

~

users

1.1	-.2	.3	.5	-.2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

items

~

.1	-.4	.2
-.5	6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-2
-1	.7	.3

Discovering features via matrix factorization

users

1		3		5		5		4	
			5		4		2	1	3
2	4		1	2	3	4	3	5	
	2	4		5		4		2	
		4	3	4	2			2	5
1		3	3		2			4	

items

{

users

1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

items

•

.1	-.4	.2
-.5	.6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-2
-1	.7	.3

{

Discovering features via matrix factorization

users

1		3		5		5		4	
		5		4		2	1	3	
2	4		1	2	3	4	3	5	
	2	4		5		4		2	
		4	3	4	2			2	5
1		3	3		2			4	

items

{

users

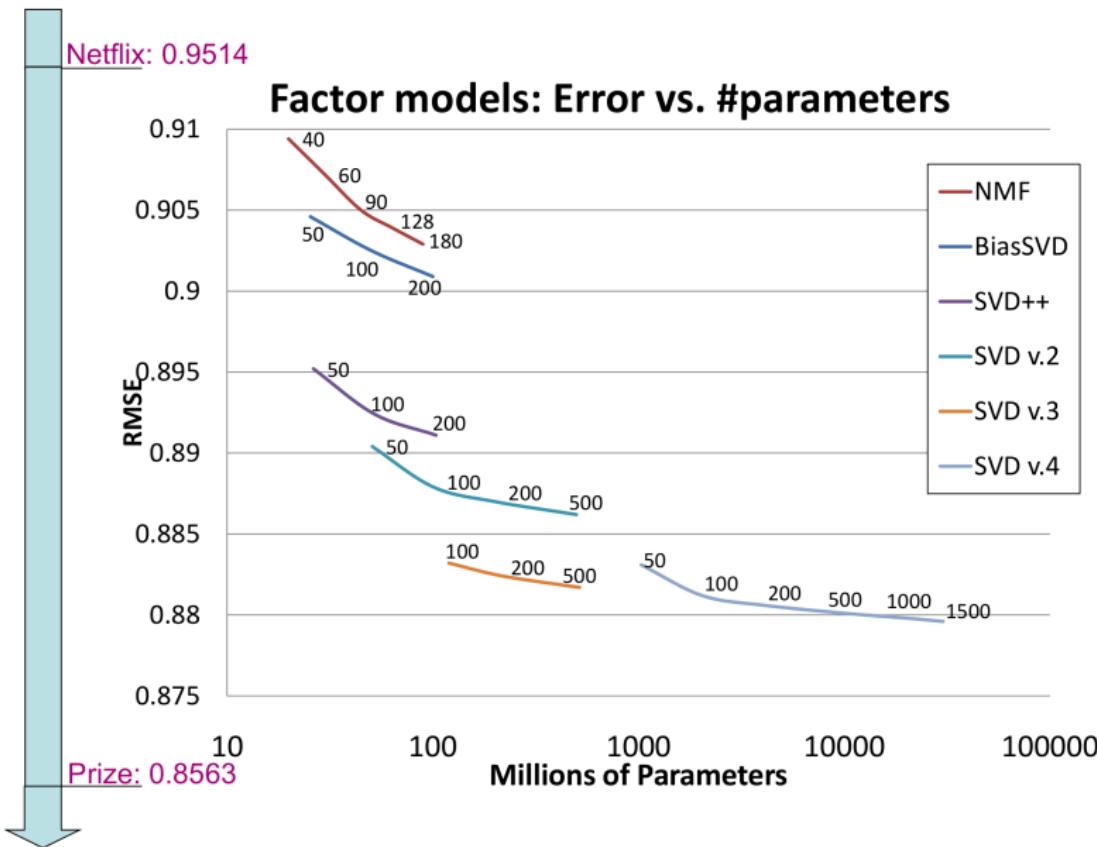
1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

{

items

.1	-.4	.2
-.5	.6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-2
-1	.7	.3

•



Finding parameters of the latent factor model

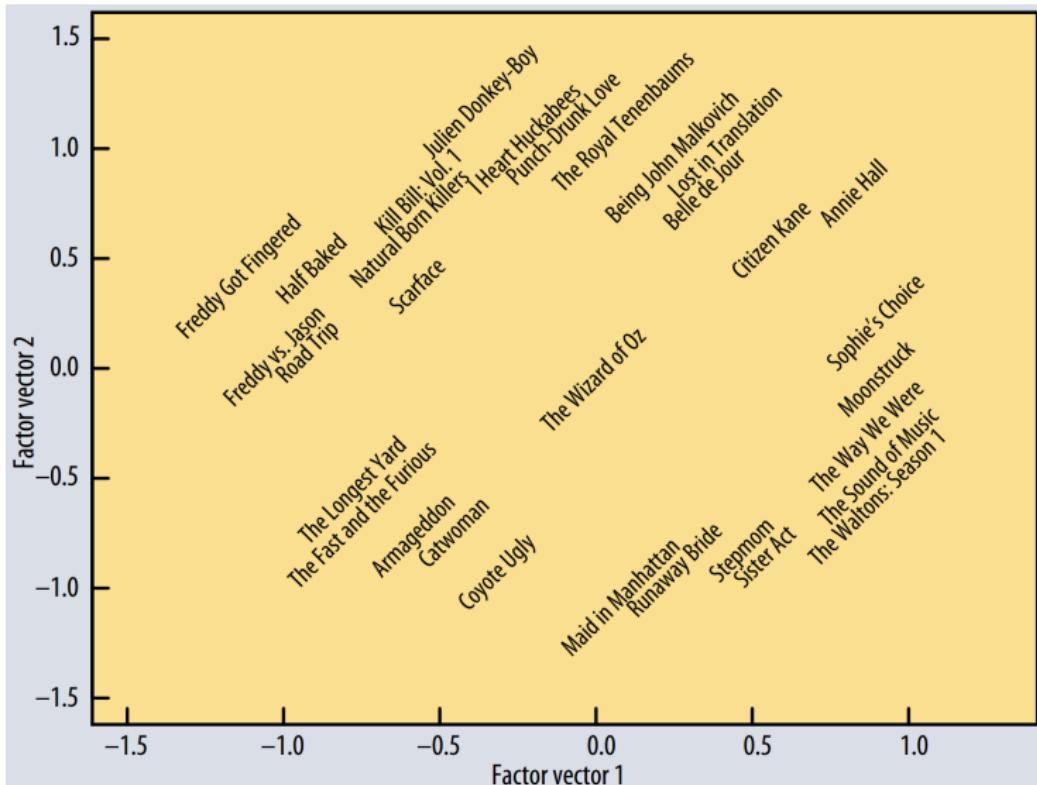
Model $f(u, i) = p_u \cdot q_i$

We will find p and q by minimizing the following objective

$$\min_{p,q} \sum_{(u,i):r_{ui}\neq?} (r_{ui} - p_u \cdot q_i)^2 + \lambda \left(\sum_{u,k} p_{u,k}^2 + \sum_{i,k} q_{i,k}^2 \right)$$

Stochastic gradient descent

Visualizing the first two factors



Visualizing the first two factors

First factor:

- ▶ left: lowbrow comedies and horror movies, aimed at a male or adolescent audience
- ▶ right: drama or comedy with serious undertones and strong female leads

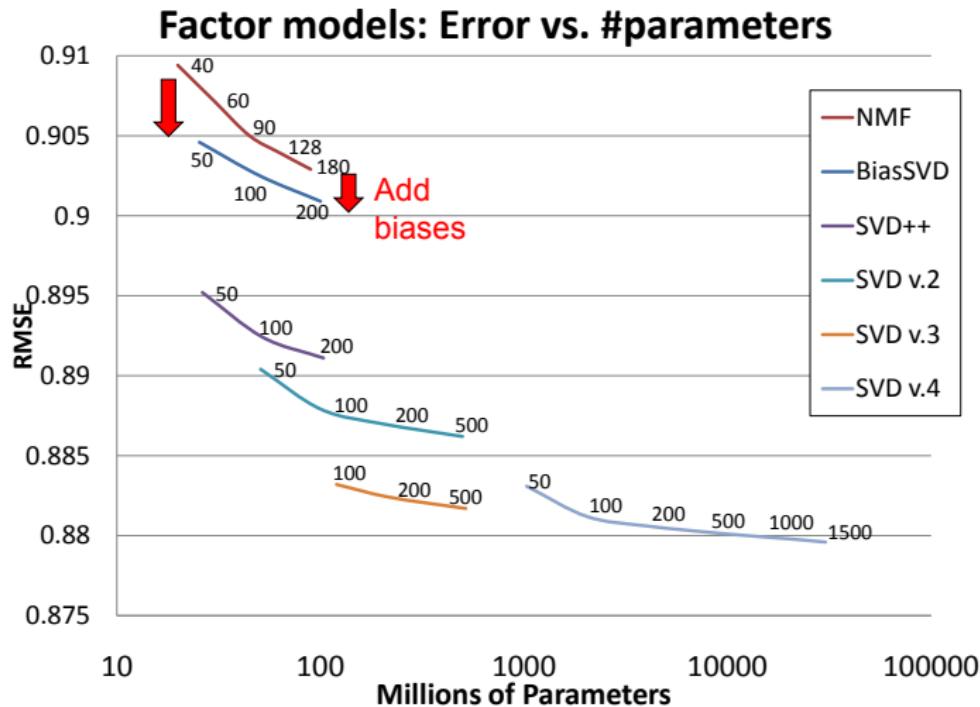
Second factor:

- ▶ top: independent, critically acclaimed, quirky films
- ▶ bottom: mainstream formulaic films

Improvements

- ▶ adding bias terms
- ▶ adding features and implicit feedback
- ▶ modelling temporal effects

Adding biases



$$f(u, i) = \mu + b_u + b_i + p_u \cdot q_i$$

Combining real and discovered features

Real features capture context

- ▶ Time of the day, what I just saw, user info, what I bought in the past

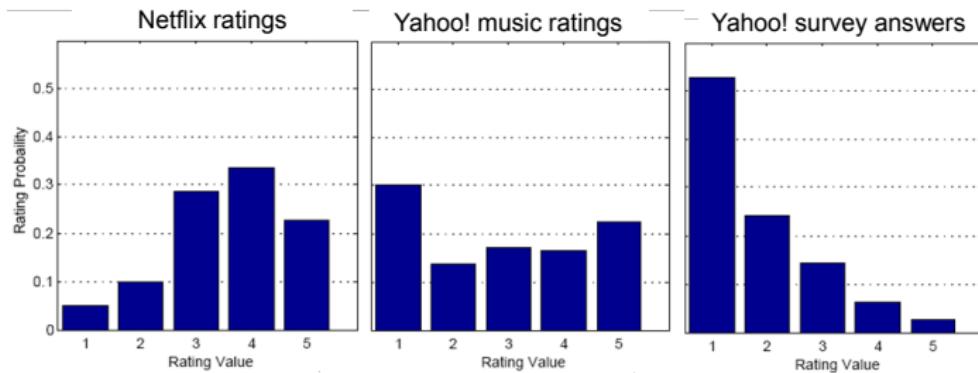
Discovered features from matrix factorization capture groups of users who behave similarly

- ▶ Users who like action movies and comedies

Mitigates cold-start problem

- ▶ Ratings for a new user from real features only
- ▶ As more information about user is discovered, matrix factorization “features” become more relevant

Implicit information



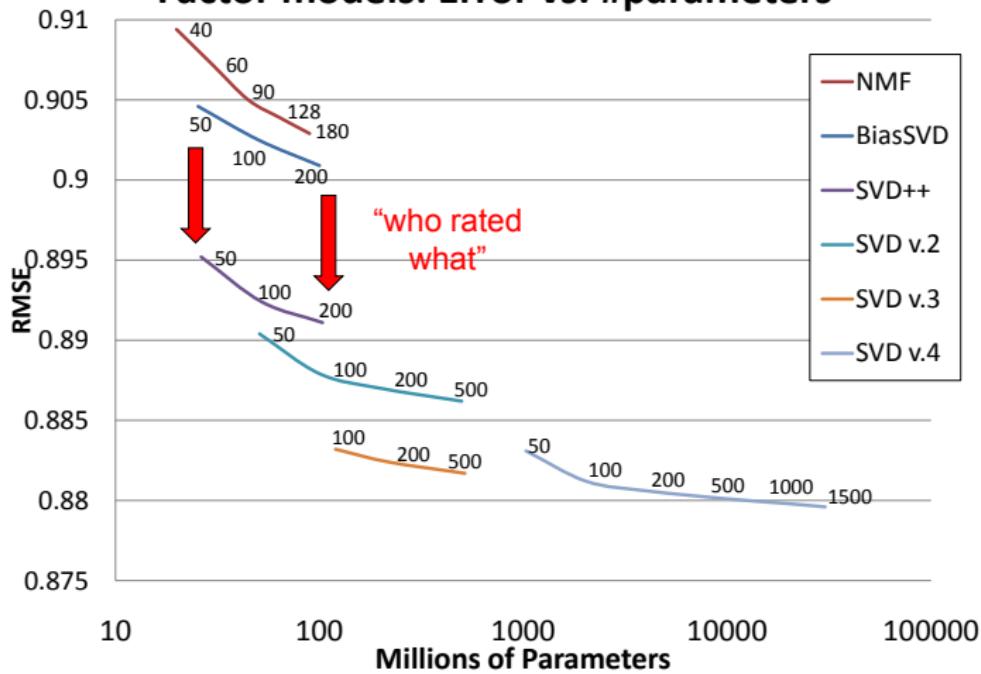
Our decision about whether to purchase a movie is a function of how we expect to rate it.

For items we have purchased, our decision to enter a rating or write a review is a function of our rating.

[http:](http://www.cs.toronto.edu/~marlin/research/papers/cfmar-uai2007.pdf)

//www.cs.toronto.edu/~marlin/research/papers/cfmar-uai2007.pdf

Factor models: Error vs. #parameters



Modeling temporal change

Time-dependent bias

Time-dependent user preferences

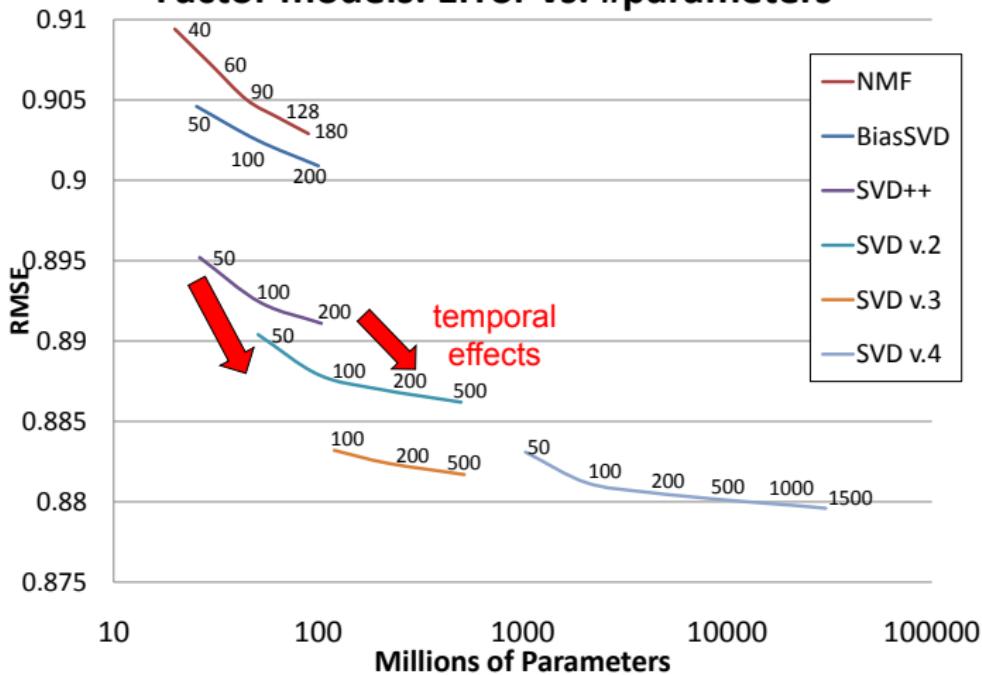
Parameterize functions b and p

$$f(u, i, t) = \mu + b_u(t) + b_i(t) + q_i \cdot p_u(t)$$

Good parametrization is the key.

<http://www.cc.gatech.edu/~zha/CSE8801/CF/kdd-fp074-koren.pdf>

Factor models: Error vs. #parameters



Moral of the story

Increasing the number of parameters does not help much, but increasing the model complexity does.

\$1,000,000 seems to be incredibly cheap to get the amount of research that was devoted to the task.

The winning solution never made it into production at Netflix.

It is not clear that a solution which changes RMSE slightly will result in hugely improved user experience.

